

Nested Monte Carlo for asian options and training

Saad Souilmi - Germain Vivier-Ardisson

March 2024

Summary

- 1 Problem statement
- 2 Monte Carlo parallelization
- 3 MLP Training
- 4 LDLt decomposition

Problem statement

$$dS_t = rS_t dt + \sigma S_t dW_t$$

The Euler-Maruyama scheme on a uniform grid $t_k = k\delta$, $k \in \{0, \dots, n\}$:

$$\forall k \in \{0, \dots, n-1\} \quad S_{t_{k+1}} = S_{t_k} \left(1 + r\delta + \sigma\sqrt{\delta}Z_{k+1} \right)$$

Where Z are i.i.d standard normal variables. Finally we compute $F(t, S_t, I_t) = e^{-r(T-t)} \mathbb{E} \left[(S_T - I_T)^+ \middle| S_t, I_t \right]$ via the proxy :

$$F(t, S_t, I_t) \approx \frac{e^{-r(T-t)}}{n_{\text{paths}}} \sum_{i=1}^{n_{\text{paths}}} \left(S_T^i - \frac{t}{T} I_t - \frac{\delta}{T} \sum_{k=1}^n S_{t+t_k}^i \right)^+$$

Parallelizing over trajectories

- We fix initial conditions (t, S_t, I_t)
- Every thread generates one sample path for following the diffusion
- To aggregate the results of all the threads we can either implement the array reduction inside the kernel or use the `numba.cuda.reduce` api.

Parallelizing over initial conditions

- The problem with the last approach is that we need to loop in cpu in order to generate multiple prices for different initial conditions.
- We can either assign a triplet (t, S_t, I_t) to each thread that will sample n_{paths} .
- Or we can assign a triplet (t, S_t, I_t) to each block and each thread will sample $\frac{n_{\text{paths}}}{n_{\text{threads per block}}}$.

We aim to learn the map :

$$F(t, S_t, I_t, T, r, \sigma) = e^{-r(T-t)} \mathbb{E} \left[(S_T - I_T)^+ \middle| S_t, I_t \right]$$

For convenience, we will write :

$$x = (t, S_t, I_t, T, r, \sigma) \text{ and } X = e^{-r(T-t)} (S_T - I_T)^+$$

We can thus write F as $F(x) = \mathbb{E}[X|x]$.

We want to find :

$$\theta^* \in \operatorname{argmin}_{\theta \in \Theta} L(\theta) = \operatorname{argmin}_{\theta \in \Theta} \mathbb{E}_{x \sim \mathcal{D}} \left[(F(x) - T_{\theta}(x))^2 \right]$$

Where \mathcal{D} is some prior distribution over the parameter space, and T_{θ} is a neural network.

$$L(\theta) = \underbrace{E_{x \sim \mathcal{D}} \left[\mathbb{E} \left[(X - T_{\theta}(x))^2 | x \right] \right]}_{=\tilde{L}(\theta)} - \mathbb{E}_{x \sim \mathcal{D}} [\mathbb{V}(X|x)]$$

$$\text{Thus } \operatorname{argmin}_{\theta \in \Theta} L(\theta) = \operatorname{argmin}_{\theta \in \Theta} \tilde{L}(\theta)$$

We can thus train our network directly on the sampled payoffs.

Parameter	Interval
S_t	$[30, 70]$
$T - t$	$[0.2, 1]$
t	$[0, 0.8]$
σ	$[0.1, 0.5]$
r	$[0, 0.1]$

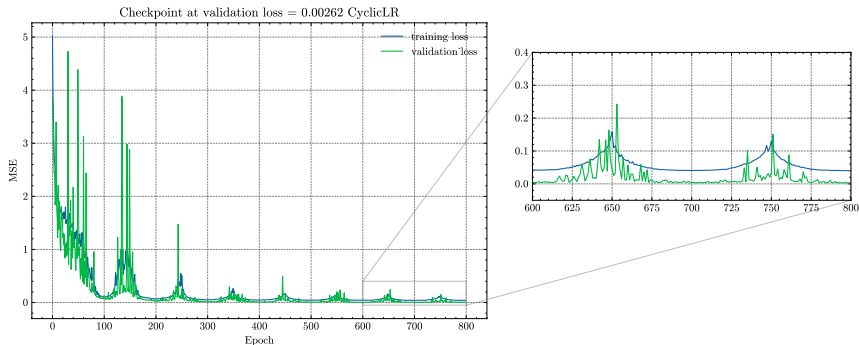
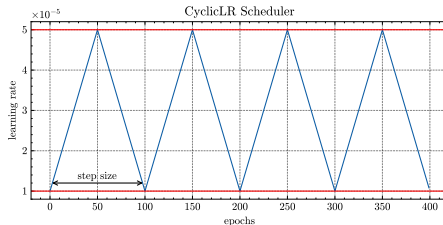
Table – Grid of parameter intervals.

- When drawing a set of parameters we draw the five parameters below with the corresponding interval, and if t is small enough, we set I_t to be equal to S_t otherwise, we sample I_t in $[0.5S_t, 2S_t]$.

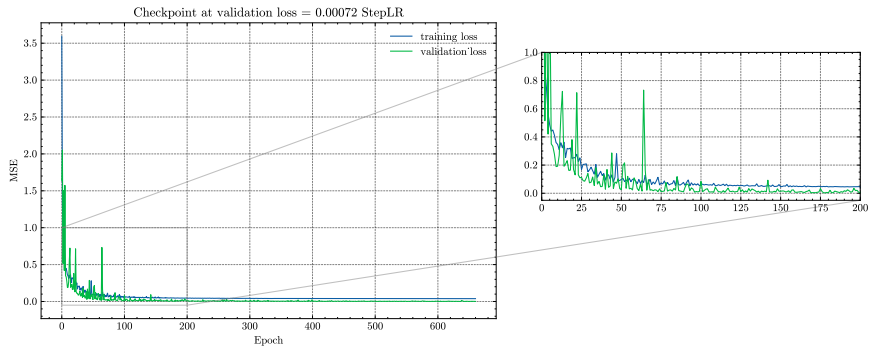
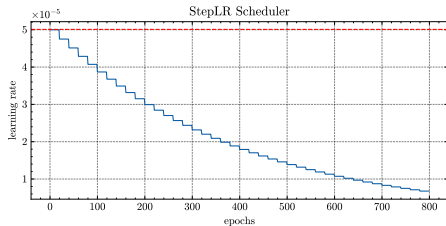
- For training data, we drew 10^6 sample parameters from the grid defined before using the scrambled Halton sequence. For each sample parameter we generated 10^3 paths.
- For validation data, we drew 10^4 sample parameters from the grid defined before using the scrambled Halton sequence. For each sample parameter we generated 10^6 paths and computed the monte carlo estimation of F .

- The neural network we trained is a fully connected MLP, with 4 hidden layers, each containing 400 neurons.
- We chose for our activation the SiLU function.
- We introduced LayerNorm between the hidden layers.
- Since we predict a positive outcome (price), we run our final output through a ReLU.

- For training we first attempted to train our network directly over payoffs.
- Training was very slow, even with big batch sizes ($\approx 3 * 10^4$), as in this case our training dataset contains 10^9 samples. And also very unstable.
- We opted to train our network on MC estimations of the price. Even though the estimations are noisy since they only utilize 1k paths, our network manages to converge.



StepLR



LDLt decomposition

- In the LDLt implementation each block works on a portion of the matrix a .
- Each block is divided into groups of n threads, and each group works on a row of the portion of the matrix associated to the block.
- $\text{int tid} = \text{threadIdx.x} \% n$ determines the column index the thread is working on.
- $\text{int Qt} = (\text{threadIdx.x} - \text{tid}) / n$ determines the index of the row that the thread is working on, within the portion of a associated with the block.
- $\text{int gb_index_x} = \text{Qt} + \text{blockIdx.x} * (\text{blockDim.x} / n)$ is the global memory index of the row the thread is working on this time in the matrix a .