

# Nested Monte Carlo for asian options and training

Saad Souilmi - Germain Vivier-Ardisson

March 2024

# Summary

- 1 Problem statement
- 2 Monte Carlo parallelization
- 3 MLP Training
- 4 LDLt decomposition

# Problem statement

$$dS_t = rS_t dt + \sigma S_t dW_t$$

The Euler-Maruyama scheme on a uniform grid  $t_k = k\delta$ ,  $k \in \{0, \dots, n\}$  :

$$\forall k \in \{0, \dots, n-1\} \quad S_{t_{k+1}} = S_{t_k} \left( 1 + r\delta + \sigma\sqrt{\delta}Z_{k+1} \right)$$

Where  $Z$  are i.i.d standard normal variables. Finally we compute  $F(t, S_t, I_t) = e^{-r(T-t)} \mathbb{E} \left[ (S_T - I_T)^+ \middle| S_t, I_t \right]$  via the proxy :

$$F(t, S_t, I_t) \approx \frac{e^{-r(T-t)}}{n_{\text{paths}}} \sum_{i=1}^{n_{\text{paths}}} \left( S_T^i - \frac{t}{T} I_t - \frac{\delta}{T} \sum_{k=1}^n S_{t+t_k}^i \right)^+$$

# Parallelizing over trajectories

- We fix initial conditions  $(t, S_t, I_t)$
- Every thread generates one sample path for following the diffusion
- To aggregate the results of all the threads we can either implement the array reduction inside the kernel or use the `numba.cuda.reduce` api.

# Parallelizing over initial conditions

- The problem with the last approach is that we need to loop in cpu in order to generate multiple prices for different initial conditions.
- We can either assign a triplet  $(t, S_t, I_t)$  to each thread that will sample  $n_{\text{paths}}$ .
- Or we can assign a triplet  $(t, S_t, I_t)$  to each block and each thread will sample  $\frac{n_{\text{paths}}}{n_{\text{threads per block}}}$ .

We aim to learn the map :

$$F(t, S_t, I_t, T, r, \sigma) = e^{-r(T-t)} \mathbb{E} \left[ (S_T - I_T)^+ \middle| S_t, I_t \right]$$

For convenience, we will write :

$$x = (t, S_t, I_t, T, r, \sigma) \text{ and } X = e^{-r(T-t)} (S_T - I_T)^+$$

We can thus write  $F$  as  $F(x) = \mathbb{E}[X|x]$ .

We want to find :

$$\theta^* \in \operatorname{argmin}_{\theta \in \Theta} L(\theta) = \operatorname{argmin}_{\theta \in \Theta} \mathbb{E}_{x \sim \mathcal{D}} \left[ (F(x) - T_{\theta}(x))^2 \right]$$

Where  $\mathcal{D}$  is some prior distribution over the parameter space, and  $T_{\theta}$  is a neural network.

$$L(\theta) = \underbrace{E_{x \sim \mathcal{D}} \left[ \mathbb{E} \left[ (X - T_{\theta}(x))^2 | x \right] \right]}_{=\tilde{L}(\theta)} - \mathbb{E}_{x \sim \mathcal{D}} [\mathbb{V}(X|x)]$$

$$\text{Thus } \operatorname{argmin}_{\theta \in \Theta} L(\theta) = \operatorname{argmin}_{\theta \in \Theta} \tilde{L}(\theta)$$

We can thus train our network directly on the sampled payoffs.

Parameter	Interval
$S_t$	$[30, 70]$
$T - t$	$[0.2, 1]$
$t$	$[0, 0.8]$
$\sigma$	$[0.1, 0.5]$
$r$	$[0, 0.1]$

Table – Grid of parameter intervals.

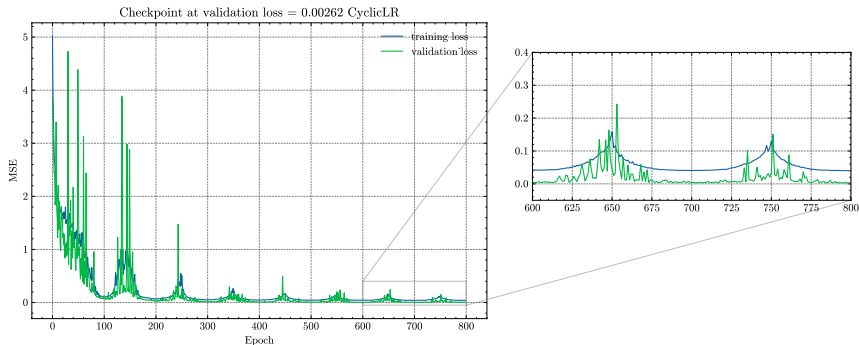
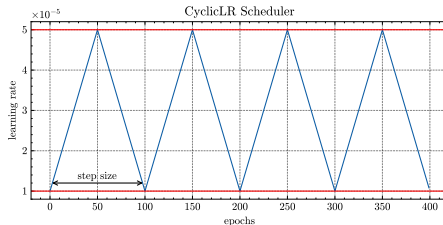
- When drawing a set of parameters we draw the five parameters below with the corresponding interval, and if  $t$  is small enough, we set  $I_t$  to be equal to  $S_t$  otherwise, we sample  $I_t$  in  $[0.5S_t, 2S_t]$ .





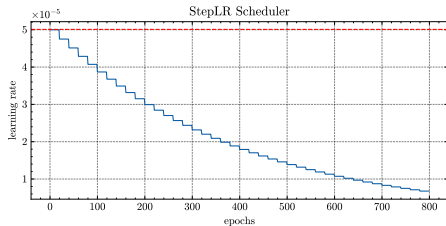
- For training data, we drew  $10^6$  sample parameters from the grid defined before using the scrambled Halton sequence. For each sample parameter we generated  $10^3$  paths.
- For validation data, we drew  $10^4$  sample parameters from the grid defined before using the scrambled Halton sequence. For each sample parameter we generated  $10^6$  paths and computed the monte carlo estimation of  $F$ .

- The neural network we trained is a fully connected MLP, with 4 hidden layers, each containing 400 neurons.
- We chose for our activation the SiLU function.
- We introduced LayerNorm between the hidden layers.
- Since we predict a positive outcome (price), we run our final output through a ReLU.

- For training we first attempted to train our network directly over payoffs.
- Training was very slow, even with big batch sizes ( $\approx 3 * 10^4$ ), as in this case our training dataset contains  $10^9$  samples. And also very unstable.
- We opted to train our network on MC estimations of the price. Even though the estimations are noisy since they only utilize 1k paths, our network manages to converge.



# StepLR



# LDLt decomposition

- In the LDLt implementation each block works on a portion of the matrix  $a$ .
- Each block is divided into groups of  $n$  threads, and each group works on a column of the portion of the matrix associated to the block.
- $\text{int tid}x = \text{threadIdx}.x \% n$  determines the row index the thread is working on.
- $\text{int Qt} = (\text{threadIdx}.x - \text{tid}x) / n$  determines the index of the column that the thread is working on, within the portion of  $a$  associated with the block.
- $\text{int gb\_index}_x = \text{Qt} + \text{blockIdx}.x * (\text{blockDim}.x / n)$  is the global memory index of the column the thread is working on this time in the matrix  $a$ .